

- (TrustCom). IEEE, 887–894.
- [34] Yuanchun Li, Hao Wen, Weijun Wang, Xiangyu Li, Yizhen Yuan, Guohong Liu, Jiacheng Liu, Wenxing Xu, Xiang Wang, Yi Sun, et al. 2024. Personal llm agents: Insights and survey about the capability, efficiency and security. *arXiv preprint arXiv:2401.05459* (2024).
 - [35] Peiyu Liu, Junming Liu, Lirong Fu, Kangjie Lu, Yifan Xia, Xuhong Zhang, Wenzhi Chen, Haiqin Weng, Shouling Ji, and Wenhai Wang. 2024. Exploring ChatGPT’s Capabilities on Vulnerability Management. In *33rd USENIX Security Symposium (USENIX Security 24)*. USENIX Association.
 - [36] Ruofan Liu, Yun Lin, Xiwen Teoh, Gongshen Liu, Zhiyong Huang, and Jin Song Dong. 2024. Less Defined Knowledge and More True Alarms: Reference-based Phishing Detection without a Pre-defined Reference List. In *33rd USENIX Security Symposium (USENIX Security 24)*. USENIX Association.
 - [37] Potsawee Manakul, Adian Liusie, and Mark JF Gales. 2023. Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models. *arXiv preprint arXiv:2303.08896* (2023).
 - [38] Dirk Merkel. 2014. Docker: lightweight linux containers for consistent development and deployment. *Linux journal* 2014, 239 (2014), 2.
 - [39] Shervin Minaee, Tomas Mikolov, Narjes Nikzad, Meysam Chenaghlu, Richard Socher, Xavier Amatriain, and Jianfeng Gao. 2024. Large language models: A survey. *arXiv preprint arXiv:2402.06196* (2024).
 - [40] Tushar Nayan, Qiming Guo, Mohammed Al Duniawi, Marcus Botacin, Selcuk Ulugac, and Ruimin Sun. 2024. {SoK}: All You Need to Know About {On-Device} {ML} Model Extraction-The Gap Between Research and Practice. In *33rd USENIX Security Symposium (USENIX Security 24)*. 5233–5250.
 - [41] Forum of Incident Response and Security Teams. 2024. Common Vulnerability Scoring System SIG. <https://www.first.org/cvss/>
 - [42] OpenAI. [n. d.]. Safety Systems. <https://openai.com/safety-systems/>.
 - [43] OpenAI. 2023. GPT-3.5: Large language model. <https://platform.openai.com>. Accessed: 2023-08-24.
 - [44] OpenAI. 2024. GPT-4o Mini: Advancing Cost-Efficient Intelligence. <https://openai.com/index/gpt-4o-mini-advancing-cost-efficient-intelligence/>.
 - [45] OpenAI. 2024. Hello GPT-4o. <https://openai.com/index/hello-gpt-4o/>.
 - [46] Walter J Price. 1989. A benchmark tutorial. *IEEE micro* 9, 5 (1989), 28–43.
 - [47] Xue Qiu, Shuguang Wang, Qiong Jia, Chunhe Xia, and Qingxin Xia. 2014. An automated method of penetration testing. In *2014 IEEE Computers, Communications and IT Applications Conference*. IEEE, 211–216.
 - [48] Elaine Rich et al. 2008. *Automata, computability and complexity: theory and applications*. Pearson Prentice Hall Upper Saddle River.
 - [49] Marcelo Invert Palma Salas and Eliane Martins. 2015. A black-box approach to detect vulnerabilities in web services using penetration testing. *IEEE Latin America Transactions* 13, 3 (2015), 707–712.
 - [50] Muzammil Shahbaz and Roland Groz. 2009. Inferring mealy machines. In *International Symposium on Formal Methods*. Springer, 207–222.
 - [51] Minghao Shao, Sofija Jancheska, Meet Udeshi, Brendan Dolan-Gavitt, Haoran Xi, Kimberly Milner, Boyuan Chen, Max Yin, Siddharth Garg, Prashanth Krishnamurthy, Farshad Khorrami, Ramesh Karri, and Muhammad Shafique. 2024. NYU CTF Dataset: A Scalable Open-Source Benchmark Dataset for Evaluating LLMs in Offensive Security. *arXiv:2406.05590* [cs.CR] <https://arxiv.org/abs/2406.05590>
 - [52] Kumar Shravan, Bansal Neha, and Bhadana Pawan. 2014. Penetration Testing: A Review. *Compusoft* 3, 4 (2014), 752.
 - [53] Ben Stock, Martin Johns, Marius Steffens, and Michael Backes. 2017. How the Web Tangled Itself: Uncovering the History of {Client-Side} Web ({In} Security)}. In *26th USENIX Security Symposium (USENIX Security 17)*. 971–987.
 - [54] The OWASP Top 10 2021 team. [n. d.]. OWASP Top 10. <https://owasp.org/Top10/>. Accessed: 2024-08-24.
 - [55] Fabian M Teichmann and Sonia R Boticiu. 2023. An overview of the benefits, challenges, and legal aspects of penetration testing and red teaming. *International Cybersecurity Law Review* 4, 4 (2023), 387–397.
 - [56] Jóakim v. Kistowski, Jeremy A Arnold, Karl Huppler, Klaus-Dieter Lange, John L Henning, and Paul Cao. 2015. How to build a benchmark. In *Proceedings of the 6th ACM/SPEC international conference on performance engineering*. 333–336.
 - [57] A Vaswani. 2017. Attention is all you need. *Advances in Neural Information Processing Systems* (2017).
 - [58] Vulhub Project. [n. d.]. Vulhub: Pre-Built Vulnerable Environments Based on Docker-Compose. <https://vulhub.org/>.
 - [59] webAI. 2024. webAI: Enterprise Grade Local AI Applications. <https://www.webai.com/>.
 - [60] Clark Weissman. 1995. Penetration testing. *Information security: An integrated collection of essays* 6 (1995), 269–296.
 - [61] Xin-Cheng Wen, Cuiyun Gao, Shuzheng Gao, Yang Xiao, and Michael R Lyu. 2024. SCALE: Constructing Structured Natural Language Comment Trees for Software Vulnerability Detection. In *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis*. 235–247.
 - [62] Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe Wang, Senjie Jin, Enyu Zhou, et al. 2023. The rise and potential of large language model based agents: A survey. *arXiv preprint arXiv:2309.07864* (2023).

- [63] Linyao Yang, Hongyang Chen, Zhao Li, Xiao Ding, and Xindong Wu. 2023. Chatgpt is not enough: Enhancing large language models with knowledge graphs for fact-aware language modeling. *arXiv preprint arXiv:2306.11489* (2023).
- [64] Mihalis Yannakakis. 1991. Testing finite state machines. In *Proceedings of the twenty-third annual ACM symposium on Theory of computing*. 476–485.
- [65] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. *arXiv:2210.03629* [cs.CL] <https://arxiv.org/abs/2210.03629>
- [66] Jiahao Yu, Xingwei Lin, Zheng Yu, and Xinyu Xing. 2024. LLM-Fuzzer: Scaling Assessment of Large Language Model Jailbreaks. In *33rd USENIX Security Symposium (USENIX Security 24)*. USENIX Association.
- [67] Xiao Yu, Lei Liu, Xing Hu, Jacky Keung, Xin Xia, and David Lo. 2024. Practitioners' Expectations on Automated Test Generation. In *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis*. 1618–1630.
- [68] Zhihang Yuan, Yuzhang Shang, Yang Zhou, Zhen Dong, Chenhao Xue, Bingzhe Wu, Zhikai Li, Qingyi Gu, Yong Jae Lee, Yan Yan, et al. 2024. Llm inference unveiled: Survey and roofline model insights. *arXiv preprint arXiv:2402.16363* (2024).
- [69] Cen Zhang, Yaowen Zheng, Mingqiang Bai, Yeting Li, Wei Ma, Xiaofei Xie, Yuekang Li, Limin Sun, and Yang Liu. 2024. How Effective Are They? Exploring Large Language Model Based Fuzz Driver Generation. In *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis*. 1223–1235.
- [70] Jianming Zhao, Wenli Shang, Ming Wan, and Peng Zeng. 2015. Penetration testing automation assessment method based on rule tree. In *2015 IEEE International Conference on Cyber Technology in Automation, Control, and Intelligent Systems (CYBER)*. IEEE, 1829–1833.
- [71] Yuchen Zhou and David Evans. 2014. SSOScan: Automated Testing of Web Applications for Single Sign-On Vulnerabilities. In *23rd USENIX Security Symposium (USENIX Security 14)*. USENIX Association, San Diego, CA, 495–510. <https://www.usenix.org/conference/usenixsecurity14/technical-sessions/presentation/zhou>