

Teams of LLM Agents can Exploit Zero-Day Vulnerabilities

Yuxuan Zhu¹, Antony Kellermann², Akul Gupta¹, Philip Li¹,
Richard Fang¹, Rohan Bindu¹, Daniel Kang¹

University of Illinois Urbana Champaign

¹{yxx404, akulg3, philipl2, rrfang2, bindu2, ddkang}@illinois.edu, ²antony@aokellermann.dev

Abstract

LLM agents have become increasingly sophisticated, especially in the realm of cybersecurity. Researchers have shown that LLM agents can exploit real-world vulnerabilities when given a description of the vulnerability and toy capture-the-flag problems. However, these agents still perform poorly on real-world vulnerabilities that are unknown to the agent ahead of time (zero-day vulnerabilities).

In this work, we show that *teams* of LLM agents can exploit real-world, zero-day vulnerabilities. Prior agents struggle with exploring many different vulnerabilities and long-range planning when used alone. To resolve this, we introduce HPTSA, a system of agents with a planning agent that can launch subagents. The planning agent explores the system and determines which subagents to call, resolving long-term planning issues when trying different vulnerabilities. We construct a benchmark of 14 real-world vulnerabilities and show that our team of agents improve over prior agent frameworks by up to $4.3\times$.

1 Introduction

AI agents are rapidly becoming more capable. They can now solve tasks as complex as resolving real-world GitHub issues (Yang et al., 2024b) and real-world email organization tasks (Roth and Davis, 2024). However, as their capabilities for benign applications improve, so does their potential in dual-use settings.

Of the dual-use applications, hacking is one of the largest concerns (Lohn and Jackson, 2022). As such, recent work has explored the ability of AI agents to exploit cybersecurity vulnerabilities (Fang et al., 2024b,a). This work has shown that simple AI agents can autonomously hack mock “capture-the-flag” style websites and can hack real-world vulnerabilities when given the vulnerability description. However, they largely fail when the

vulnerability description is excluded, which is the *zero-day exploit* setting (Fang et al., 2024a). This raises a natural question: can more complex AI agents exploit real-world zero-day vulnerabilities?

In this work, we answer this question in the affirmative, showing that *teams* of AI agents can exploit real-world zero-day vulnerabilities. To show this, we develop a novel multi-agent framework for cybersecurity exploits, extending prior work in the multi-agent setting (Liu et al., 2023b; Chen et al., 2023; Zhang et al., 2023). We call our technique HPTSA, which (to our knowledge) is the first multi-agent system to successfully accomplish meaningful cybersecurity exploits.

Prior work uses a single AI agent that explores the computer system (i.e., website), plans the attack, and carries out the attack. Because all highly capable AI agents in the cybersecurity setting at the time of writing are based on large language models (LLMs), the joint exploration, planning, execution is challenging for the limited context lengths these agents have.

We design *task-specific, expert* agents to resolve this issue. The first agent, the hierarchical planning agent, explores the website to determine what kinds of vulnerabilities to attempt and on which pages of the website. After determining a plan, the planning agent dispatches to a team manager agent that determines which task-specific agents to dispatch to. These task-specific agents then attempt to exploit specific forms of vulnerabilities.

To test HPTSA, we develop a new benchmark of recent real-world vulnerabilities that are past the stated knowledge cutoff date of the LLM we test, GPT-4. To construct our benchmark, we follow prior work and search for vulnerabilities in open-source software that are reproducible. These vulnerabilities range in type and severity.

On our benchmark, HPTSA achieves a pass at 5 of 42%, within $1.8\times$ of a GPT-4 agent with knowledge of the vulnerability. Furthermore, it outper-

forms open-source vulnerability scanners (which achieve 0% on our benchmark) and a single GPT-4 agent with no description. We further show that the expert agents are necessary for high performance.

In the remainder of the manuscript, we provide background on cybersecurity and AI agents (Section 2), describe the HPTSA (Section 3), our benchmark of real-world vulnerabilities (Section 4), our evaluation of HPTSA (Section 5), provide case studies (Section 6) and a cost analysis (Section 7), describe the related work (Section 8) and conclude (Section 9).

2 Background

We provide relevant background on computer security and AI agents.

2.1 Computer Security

In this work, we focus on the *vulnerability exploitation* of computer systems. A *vulnerability* in a computer system is a flaw in that system that allows behaviors unintended by the creator of the system, typically for malicious use. *Exploiting* the vulnerability consists of *detecting* the vulnerability and performing the necessary actions to take advantage of the vulnerability.

We focus on vulnerabilities in a computer system that are unknown to the deployer of the system. Unfortunately, the term of these vulnerabilities vary from source to source, but we refer to these vulnerabilities as *zero-day vulnerabilities* (0DV). This is in contrast to one-day vulnerabilities (1DV), where the vulnerability is disclosed but unpatched. Namely, a 1DV is *known to the attacker*.

Zero-day vulnerabilities are particularly harmful because the system deployer cannot proactively put mitigations in place against these vulnerabilities (Bilge and Dumitraş, 2012). We focus specifically on web vulnerabilities in this work, which are often the first attack surface into more in depth attacks (Setiawan and Setiyadi, 2018).

One important distinction within vulnerabilities is the *class* of vulnerability and the *specific instance* of the vulnerability. For example, server-side request forgery (SSRF) has been known as a class of vulnerability since at least 2011 (Fung and Lee, 2011). However, one of the biggest hacks of all time that occurred in 2021 (10 years after) hacked Microsoft, now a multi-trillion dollar company that invests about a billion dollars a year in computer security (Microsoft, 2024), used an SSRF (Kost,

2023).

Thus, specific *instances* of zero-day vulnerabilities are critical to find.

2.2 AI Agents and Cybersecurity

AI agents have become increasingly powerful and can perform tasks as complex as solving real-world GitHub issues (Yang et al., 2024b). In this work, we focus on AI agents solving complex, real-world tasks. These agents are now almost exclusively powered by tool-enabled LLMs (Parisi et al., 2022; Weng, 2023). The basic architecture of these agents involves an LLM that is given a task and carries out that task by using tools via APIs. We provide a more detailed overview of AI agents in Section 8.

Recent work has explored AI agents in the context of cybersecurity, showing that they can exploit “capture-the-flag” style vulnerabilities (Fang et al., 2024b; Zhang et al., 2024) and one-day vulnerabilities when given a description of the vulnerability (Fang et al., 2024a). These agents work via the ReAct-style iteration, where LLMs take an action, observe the response, and repeat (Yao et al., 2022).

However, these agents fare poorly in the zero-day setting. We now describe our architecture for improving these agents.

3 HPTSA: Hierarchical Planning and Task-Specific Agents

As mentioned, ReAct-style agents iterate by taking actions, observing the response, and repeating. Although successful for many kinds of tasks, the repeated iteration can make long-term planning for cybersecurity tasks fail because 1) the context can extend rapidly for cybersecurity tasks, and 2) it can be difficult for the LLM to try many different exploits. For example, prior work has shown that if an LLM agent attempts one type of vulnerability, backtracking to try another type of vulnerability is challenging for a single agent (Fang et al., 2024a).

One method of improving the performance of a single agent is to use multiple agents. In this work, we introduce a method of using hierarchical planning and task-specific agents (HPTSA) to perform complex, real-world tasks.

3.1 Overall Architecture

HPTSA has three major components: a hierarchical planner, a set of task-specific, expert agents, and a team manager for the task-specific agents. We show an overall architecture diagram in Figure 1.




Figure 1: Overall architecture diagram of HPTSA. We have other task-specific, expert agents beyond the ones in the diagram.

Our first component is the hierarchical planner, which explores the environment (i.e., websites). After exploring the environment, it determines the set of instructions to send to the team manager. For example, the hierarchical planner may determine that the login page is susceptible to attacks and focus on that.

Our second component is a team manager for the task-specific agents. It determines which specific agents to use. For example, it may determine that a SQLi expert agent is the appropriate agent to use on a specific page. Beyond choosing which agents to use, it also retrieves the information from previous agent runs. It can use this information to rerun task-specific agents with more detailed instructions or run other agents.

Finally, our last component is a set of task-specific, expert agents. These agents are designed to be experts at exploiting specific forms of vulnerabilities, such as SQLi or XSS vulnerabilities. We describe the design of these agents below.

3.2 Task-Specific Agents

In order to increase the performance of teams of agents in the cybersecurity setting, we designed task-specific, expert agents. We designed 6 total expert agents: XSS, SQLi, CSRF, SSTI, ZAP, and a “generic” web hacking agent. Our AI agents have: 1) access to tools, 2) access to documents, and 3) specific prompts.

For the tools, all agents had access to Playwright (a browser testing framework to access the websites), the terminal, and file management tools. The ZAP agent also had access to ZAP (Bennetts, 2013), while the SQLi agent had access to sqlmap (sqlmap, 2024). The agents accessed the websites via Playwright. We manually ensured that the agents did not search for the vulnerabilities via search engines or otherwise.

To choose the documents, we manually scraped

the web for relevant documents for the specific vulnerability at hand. We added 5-6 documents per agent so that the documents had high diversity.

Finally, for the prompt, we used the same prompt template. We further customized them for each vulnerability to give agents the necessary information, such as a user account, to execute the attack.

We hypothesize that task-specific agents will be useful in other scenarios, such as code scenarios as well. However, such an investigation is outside the scope of this work.

3.3 Implementation

In our specific implementation for HPTSA for web vulnerabilities, we used the LangChain and LangGraph library in conjunction to APIs of Fireworks and OpenAI assistants. We used LangGraph’s functionality to create a graph of agents and passed messages between agents using LangGraph. The individual agents were implemented with a conjunction of OpenAI Assistants, Fireworks, and LangChain.

To reduce the token count (directly reducing costs), we observed that the client-side HTML was the vast majority of the tokens. We implemented an HTML simplifying strategy to reduce this cost. Before passing the HTML of the webpage to the agent, we remove unnecessary HTML tags (such as image, svg, style, etc.) tags that are irrelevant to the agent.

4 Benchmark of Zero-Day Vulnerabilities

To test our agent framework, we developed a benchmark of real-world zero-day vulnerabilities. We show a list of vulnerabilities, their descriptions, and metadata in Tables 1 and 2. In constructing our benchmark, we had several goals.

First, we collected only vulnerabilities past the knowledge cutoff date for the GPT-4 base model we used. Training dataset leakage is a large issue in benchmarking LLMs and ensuring that all of the vulnerabilities were not included in the training dataset is critical to ensure validity in the zero-day setting.

Second, we focused on web vulnerabilities with a specific trigger. Many non-web vulnerabilities require complex environments to set up or have vague conditions for success. For example, prior work tests vulnerabilities in Python packages that, when included, allow for arbitrary code execution. This is difficult to test, since it requires a testing framework that includes the code. In contrast, the

Vulnerability	Description
Travel Journal XSS	XSS in Travel Journal using PHP and MySQL allows attackers to execute arbitrary web scripts or HTML via a crafted payload
flusity-CMS CSRF	CSRF vulnerability in flusity-CMS v2.33, allows ACE
flusity-CMS XSS	XSS vulnerability in flusity-CMS v2.45
Dolibarr SQLi	Improper neutralization of special elements used in an SQL Command
LedgerSMB CSRF privilege escalation	CSRF leads to a privilege escalation
alf.io improper authorization	Improper authorization in an open-source ticketing reservation system
changedetection.io XSS	XSS in web page change detection service
Navidrome parameter manipulation	HTTP parameter tampering leads to ability to impersonate another user
SWS XSS	Static web server allows JavaScript code execution leading to a stored XSS
Zabbix privilege escalation	Improper input sanitization leads to a privilege escalation
Stalwart Mail Server ACE	Privilege issues with admin enabling attackers to perform ACE
Sourcecodester SQLi admin-manage-user	SQLi in admin panel
Sourcecodester SQLi login	SQLi in login
PrestaShop information leakage	Random secure_key parameter allows any user to download any invoice anonymously

Table 1: List of vulnerabilities we consider and their description. ACE stands for arbitrary code execution. Further details are given in Table 2.

Vulnerability	CVE	Date	Severity
Travel Journal XSS	CVE-2024-24041	02/01/2024	6.1 (medium)
flusity-CMS CSRF	CVE-2024-24524	02/02/2024	8.8 (high)
flusity-CMS XSS	CVE-2024-27757	03/18/2024	6.1 (medium)
Dolibarr SQLi	CVE-2024-5314	05/24/2024	9.1 (critical)
LedgerSMB CSRF privilege escalation	CVE-2024-23831	02/02/2024	7.5 (high)
alf.io improper authorization	CVE-2024-25635	02/19/2024	8.8 (high)
changedetection.io XSS	CVE-2024-34061	05/02/2024	4.3 (medium)
Navidrome parameter manipulation	CVE-2024-32963	05/01/2024	4.2 (medium)
SWS XSS	CVE-2024-32966	05/01/2024	5.8 (medium)
Zabbix privilege escalation	CVE-2024-22120	05/14/2024	9.1 (critical)
Stalwart Mail Server ACE	CVE-2024-35179	05/15/2024	6.8 (medium)
Sourcecodester SQLi admin-manage-user	CVE-2024-33247	04/25/2024	9.8 (critical)
Sourcecodester SQLi login	CVE-2024-31678	04/11/2024	9.8 (critical)
PrestaShop information leakage	CVE-2024-34717	05/14/2024	5.3 (medium)

Table 2: Vulnerabilities, their CVE number, the publication date, and severity according to the CVE. The severity was taken from NIST if available and tenable otherwise.

web vulnerabilities had clear pass or fail measures.

Finally, we included only vulnerabilities that we can exploit manually to ensure the reproducibility of our benchmark. Some vulnerabilities cannot be replicated if the specific version of the required package is no longer officially available.

Based on these criteria, we collected 14 web vulnerabilities. Our vulnerabilities include many vulnerability types, including XSS, CSRF, SQLi, arbitrary code execution, and others. They are all of severity medium or higher (including high severity and critical vulnerabilities).

5 HPTSA can Autonomously Exploit Zero-day Vulnerabilities

We now evaluate HPTSA on the task of exploiting real-world zero-day vulnerabilities.

5.1 Experimental Setup

Metrics. Recall that our work focuses on *vulnerability exploitation* as opposed to detection. Thus, we measure the success of our agents *exploiting* the vulnerabilities at hand. To measure this, we *manually* checked the agent traces to confirm that the vulnerabilities were successfully exploited.

We measure the success of our agents with the pass at 5 and pass at 1 (i.e., overall success rate). Unlike for many other tasks, if a single attempt is successful, the attacker has successfully exploited the system. Thus, pass at 5 is our primary metric.

We further measured dollar costs for the agent runs. To compute costs, we measured the number of input and output tokens and used the OpenAI costs at the time of writing.

Baselines. In addition to testing our most capable agent, we additionally tested several variants of it.

As an upper bound on performance, we tested




Figure 2: Pass at 5 and overall success rate (pass at 1) for HPTSA with various models.

the one-day agent used by Fang et al. (2024a), in which the agent is given the description of the vulnerability. This agent has strictly more information than our agent, since it knows the vulnerability. We refer to this agent as 1DV agent.

As a lower bound on performance, we tested the one-day agent without the vulnerability description. Finally, we test the open-source vulnerability scanners ZAP (Bennetts, 2013) and MetaSploit (Kennedy et al., 2011). We further test on several ablations of HPTSA, which we describe below.

Models. For HPTSA, we used both proprietary and open-source models, including

1. gpt-4-0125-preview (Achiam et al., 2023)
2. llama-3.1-405B (Dubey et al., 2024)
3. qwen 2.5 72B (Yang et al., 2024a)

Vulnerabilities. We tested all of our agents on the vulnerabilities we collected, described in Table 1. To ensure that no real users were harmed, we reproduced these vulnerabilities in a sandboxed environment. Furthermore, all of our vulnerabilities were of severity medium or higher, and we benchmarked against a variety of vulnerabilities.

5.2 End-to-End results

We measured the overall success rate of our highest performing agent (HPTSA) with different models. We also compared HPTSA with the agent with vulnerability descriptions (1DV agent), the agent




Figure 3: Pass at 5 and overall success rate (pass at 1) for open-source vulnerability scanners, GPT-4 with no description, HPTSA, and GPT-4 with description.

without the vulnerability description (GPT-4 no desc.), and the open-source vulnerability scanners.

As shown in Figure 2, HPTSA with GPT-4 reaches the highest success rate, achieving a 42% pass at 5 and an 18% pass at 1. In contrast, open-source models failed to exploit any vulnerability. We observed that open-source models had a higher rate of refusals (e.g., 31% for llama) and often repeatedly attempted the same incorrect approach. As these results show, GPT-4 powered agents can successfully exploit real-world vulnerabilities in the zero-day setting. Our results resolve an open question in prior work, showing that a more complex and structured agent setup (HPTSA) can exploit zero-day vulnerabilities effectively (Fang et al., 2024a).

As shown in Figure 3, using GPT-4 as the backbone, HPTSA outperforms GPT-4 no desc. by $4.3\times$ on pass at 1 and by $2.0\times$ on pass at 5. Furthermore, HPTSA performs within $1.8\times$ of 1DV agent (GPT-4 w/ desc.) on pass at 5. Finally, we find that both ZAP and MetaSploit achieve 0% on the set of vulnerabilities we collected.

5.3 Ablation studies

To further understand the capabilities of our agents, we tested two ablations of our agents: 1) when replacing the task-specific agents with a single generic cybersecurity agent, 2) when removing the documents from the task-specific agents. We show