




Fig. 4. High-Level Architecture Diagram of our Prototype *cochise*. The **PLANNER** is responsible for creating the high-level task plan with tasks to execute while the **EXECUTOR** is responsible for executing those tasks. Information (the to be achieved task and its description) flows from the **PLANNER** to the **EXECUTOR**, the **EXECUTOR** returns a summary of its results and a collected shell history.

```

1 1. Perform basic network enumeration on 192.168.56.0/24, excluding 192.168.56.1 and 192.168.56.107
2 1.1. Identify which hosts are Windows systems and potential domain controllers
3
4 2. Enumerate domain accounts using netexec, focusing on discovered Windows hosts that appear to be
5 domain controllers or file servers

```

Fig. 5. Example of an initially generated **PLANNER** State (Pentest-Task-Tree) by OpenAI's o1-GPT-4o. As the LLM does only possess limited information about the target environment, the generated PTT only consists of initial network enumeration tasks.

The prototype has a **EXECUTOR** round limit of 10 rounds. After this limit has been reached, the **EXECUTOR** is stopped and the LLM instructed to create a final summary of the **EXECUTOR** run. If the **EXECUTOR** was able to finish the task before the 10 round limit has been reached, the summary is created during the last **EXECUTOR** round.

```

1 ...
2 3. Perform offline password cracking on discovered Kerberos hash
3 3.1. Use a cracking tool (john or hashcat) with /usr/share/wordlists/rockyou.txt to attempt
4 cracking missandei@ESSOS.LOCAL's hash.
5 3.2. If successful, validate the credentials with netexec to confirm domain access.
6 3.2.1. Findings:
7 - The provided hash for missandei@ESSOS.LOCAL could not be loaded by john or hashcat
8 - Both tools reported issues with the hash format
9 3.3. Re-verify Kerberos hash format for Missandei@ESSOS.LOCAL
10 3.3.1. Confirm the correct format for the $krb5asrep$ hash output from impacket-GetNPUsers.
11 3.3.2. If needed, re-capture missandei's Kerberos ASREP hash in a recognized format that john
12 or hashcat can handle.
13 3.3.3. Attempt offline cracking again with the corrected hash format.
14 3.3.4. If cracking is successful, proceed to confirm credentials with netexec.
15 3.3.5. Findings:
16 - Successfully retrieved Missandei's ASREP hash in a compatible format
17 - Cracked password is "fr3edom"
18 - Verified valid domain credentials (essos.local\missandei:fr3edom)
19 ...

```

Fig. 6. Excerpt of the PLANNER state (PTT) generated by OpenAI's o1-GPT-4o after 10 *update-strategy* rounds (see Section B.2 for full listing). The PLANNER was able to integrate findings, e.g., the retrieved and cracked hash of *missandei*, within the PTT. Please note, that originally (3.2.1) the LLM was not able to successfully crack the password hash. It re-captured the network hash (3.3.2) and re-performed the hash cracking operation (3.3.3) and was able to retrieve the plain-text password shown in 3.3.5. This is an example of a successful multi-step attack also including auto-repair (see Section 6.4.3).

```

1 Task: 1.1. Perform an nmap scan on 192.168.56.0/24 (excluding 192.168.56.1 192.168.56.100
      192.168.56.107) using only eth1 to identify which hosts are accessible and what ports are open.
2
3 Context: This will help determine the live hosts and key services running within the target network
      prior to attempting user or service-based attacks. We have not yet collected any specific
      credentials or identified individual accounts, so our focus is to gather network-based
      information first.

```

Fig. 7. Example Task and Context generated by the PLANNER for the EXECUTOR. The PLANNER generates these based on its current knowledge of the target environment (as documented in the PTT). The task and context should incorporate enough information for the EXECUTOR to successfully finish the given task. This example was generated by OpenAI's o1 for the initial network enumeration and thus contains only limited target information.

4.3 Interactions between PLANNER and EXECUTOR

The EXECUTOR returns the executed task, an executive summary, and a list of all executed commands and their outputs back to the PLANNER. This data, together with the existing PTT, is subsequently used by the PLANNER to update its PTT.

The EXECUTOR itself stores no local information, i.e., its history of executed commands and their result is cleared after each EXECUTOR run. This mandates that the PLANNER has to integrate all relevant pen-test state information within its PTT. An elegant result of this design decision is that by starting the PLANNER with a stored updated PTT, an old penetration-test run can be resumed.

We explicitly wanted the PLANNER to gain as much information as possible, i.e., include both the EXECUTOR's summary as well as its raw data in form of the executed commands and its output. We accept that this will result in higher prompting costs, esp. when using the expensive o1 reasoning model for the PLANNER. But we decided to refrain from premature optimization and focus on better understanding the PLANNER's behavior before reducing costs. As our

```

1 Task: 3.1 - Use netexec (nxc) to attempt authentication with default passwords against the discovered
  domain controllers
2
3 Context: Potential next tasks include: (2.2.1) finalize the custom user list by adding the OSINT
  usernames from /root/osint_users.txt, or (3.1) validate discovered usernames via netexec (nxc).
  The most direct step toward confirming valid accounts-and thus moving closer to domain dominance-
  is (3.1). We already have many Game of Thrones-inspired domain usernames from /root/osint_users.
  txt (e.g., 'j.snow', 'd.targaryen', etc.). We will target the domain controllers at 192.168.56.10
  (sevenkingdoms.local), 192.168.56.11 (winterfell.north.sevenkingdoms.local), and 192.168.56.12 (
  essos.local) over SMB/WinRM.
4
5 We can try a small set of common default passwords (e.g., 'Password1', 'Winter2022', 'Welcome1')
  against a subset of the discovered usernames to avoid lockouts. The netexec command format will
  look like:
6
7 nxc smb 192.168.56.10 -u <username1> <username2> -p <candidate_password> --port 445 --interface eth1
8
9 (Adjust the target IP among the three domain controllers, and test only a few usernames/passwords at
  once to minimize lockout risk.)

```

Fig. 8. Example Task and Context generated by the PLANNER for the EXECUTOR. The PLANNER generates these based on its current knowledge of the target environment (as documented in the PTT). The task and context should incorporate enough information for the EXECUTOR to successfully finish the given task. This example was generated by OpenAI’s o1 and tasked the EXECUTOR to perform a targeted password-spraying attack. Note that the PLANNER included scenario-specific information and provided detailed information about the target environment that it gathered during prior execution steps.

results will show, running our prototype occurs at substantially lower costs than employing a professional pen-tester does, so we believe that this decision is acceptable.

We added a monetary fail-safe though: if the size of the passed command history is larger than 100000 bytes, the command line history is removed from the PLANNER call. The PLANNER thus depends only on the EXECUTOR’s summary. We utilized *LangChain’s langchain_core.messages.utils.trim_message* to fit the shell history into the EXECUTOR LLM’s context size.

5 Evaluation

We performed our evaluation according to our Experiment Design (Section 3). Tables 3, 4, 5, 6, and 7 overview our quantitative results for each evaluated LLM.

Performed Rounds describes the workload distribution within the prototype. A PLANNER round occurs every time the high-level PLANNER updates its PTT and selects a new task to be executed by the EXECUTOR. While the EXECUTOR tries to achieve the delegated task, EXECUTOR rounds occur. During the EXECUTOR’s lifetime, multiple *Commands* can be issued. The amount of EXECUTOR calls and executed system commands within an EXECUTOR round can differ: the EXECUTOR optionally performs an additional LLM call after an unsuccessful run to create a summary, and during each EXECUTOR round multiple system commands can be issued to be executed in parallel.

As described in Section 3.5 we utilize human penetration-testers to evaluate the efficacy of LLMs. They analyze the provided execution traces and note compromised user accounts (*done*), well-chosen attacks that failed but were on the right track (*Almost*), and promising *leads* written down by the LLM into the PTT but not being followed up on. Comprised user accounts (*Done*) were identified by their respective known password. Compromised accounts rated *Almost* must be attacks that targeted a relevant attack vector but failed due to a minimal problem, e.g., performing

Table 3. Overview of GPT-4o’s run results.

Run	Performed Rounds			Results			Tokens Planner		Tokens Executor		Cost	
	PLANNER	EXECUTOR	Commands	Done	Almost	Lead	Prompt	Compl.	Prompt	Compl.	Cost	per User
run-20250516-113002	49	4.31 ± 2.77	3.78 ± 2.87	2	3	6	544.56	190.4	956.94	25.78	\$4.81	\$2.41
run-20250516-140100	32	4.38 ± 2.34	4.56 ± 3.34	0	3	3	243.67	59.59	293.73	19.30	\$1.76	
run-20250516-161010	37	4.38 ± 2.78	4.14 ± 3.00	0	2	4	405.5	139.42	374.81	39.99	\$3.17	
run-20250516-181043	27	3.41 ± 2.29	3.15 ± 3.56	0	1	1	216.1	48.65	195.35	109.59	\$2.39	
run-20250517-102109	21	4.14 ± 2.56	4.57 ± 5.68	0	1	4	171.03	33.11	395.38	14.38	\$1.56	
run-20250517-173859	35	3.57 ± 2.16	3.69 ± 2.75	0	1	3	275.31	70.06	262.29	18.73	\$1.89	
Average	33.5	4.06	3.95	0.33	1.83	3.50	309.36	90.21	413.08	37.96	\$2.59	\$ 2.41
		± 2.52	± 3.42				± 139.91	± 61.31	± 276.39	± 36.22	± \$1.23	

Executed Commands are summarized per PLANNER-Round. Within results, *done* designates fully compromised user accounts, *almost* attacks that failed due to a minimal error, and leads are designated as concrete vulnerabilities that the PLANNER has included within the PTT to follow-up to (detailed in Section 3.5. All Token costs are given in kilo-Tokens (kTokens).

Table 4. Overview of DEEPSEEK-V3’s run results.

Run	Performed Rounds			Results			Tokens Planner		Tokens Executor		Cost	
	PLANNER	EXECUTOR	Commands	Done	Almost	Lead	Prompt	Compl.	Prompt	Compl.	Cost	per User
run-20250522-113839	22	2.73 ± 1.86	2.91 ± 2.22	0	3	3	275.01	100.16	134.22	10.71	\$ 0.17	
run-20250522-134507	40	3.15 ± 2.32	3.02 ± 3.21	1	2	3	405.41	120.26	440.32	24.15	\$ 0.27	\$ 0.27
run-20250522-164357	20	4.10 ± 2.49	3.3 ± 2.72	0	4	3	223.84	63.46	308.17	15.12	\$ 0.16	
run-20250522-184230	29	2.79 ± 1.92	2.17 ± 2.16	1	1	4	362.83	132.53	318.09	13.36	\$ 0.25	\$ 0.25
run-20250522-204757	27	3.26 ± 2.40	3.52 ± 2.81	0	2	2	295.75	92.39	298.09	17.54	\$ 0.21	
run-20250523-122103	20	3.35 ± 1.87	2.35 ± 1.87	0	2	3	208.20	74.33	134.88	11.12	\$ 0.13	
Average	26.33	3.19	2.89	0.33	2.33	3.00	295.17	97.19	272.3	15.33	\$ 0.20	\$ 0.26
		± 2.18	± 2.63				± 77.19	± 26.36	± 118.51	± 5.01	± \$ 0.06	

Executed Commands are summarized per PLANNER-Round. Within results, *done* designates fully compromised user accounts, *almost* attacks that failed due to a minimal error, and leads are designated as concrete vulnerabilities that the PLANNER has included within the PTT to follow-up to (detailed in Section 3.5. All Token costs are given in kilo-Tokens (kTokens).

password-spraying with a scenario-specific password list but not including the right password. *Leads* occurred when concrete evidence of a potential exploitable vulnerability was added to the PTT.

We include used token counts, separated into input and output tokens, within the tables, and calculated the occurring cost of each run based upon the token counts. The single exception is QWEN3: as this model was hosted on a rented virtual machine, we sum the duration of LLM-calls to calculate the hosting cost that occurred during LLM execution.

The included timestamps allow to match our findings to the raw log traces provided within our public github project repository²².

5.1 Non-Reasoning LLMs: OpenAI GPT-4o and DEEPSEEK-V3

We start our evaluation with “traditional” non-reasoning LLMs which were the mainstay of used LLMs between 2023–2025²³. We are analyzing both a closed-weight model (OpenAI’s GPT-4o) as well as an open-weight model (DeepSeek’s DEEPSEEK-V3). The latter is capable of running on-premise given sufficient hardware.

²²<https://github.com/andreashappe/cochise>

²³OpenAI’s made ChatGPT publicly available in November 2022; its o1-preview reasoning model was made generally available in December 2024.




Fig. 9. Attack Vectors pursued by the different LLMs. For each attack vector we detail the percentage of runs in which the respective attack vector was included. The attack must have been of sufficient quality, i.e., it must have been parametrized to our target environment (fitting the *Almost-There* and *Done* categories from our per-LLM overview tables). Please note, that QWEN3’s result were a result of QWEN3 not being able to integrate findings into its PTT and thus re-iterating over the initial network/service scanning steps. Results indicate that non-reasoning LLMs (DEEPSEEK-V3 and GPT-4o) possess sufficient penetration-testing knowledge to perform attacks while reasoning LLMs (GEMINI-2.5-FLASH and the combination of OpenAI’s o1 and GPT-4o) increase the consistency of performed attacks.

5.1.1 Comparison between GPT-4o and DeepSeek-V3. Both models were not able to routinely compromise user accounts (0.33 compromised user accounts per 2 hours). The amount of *compromised* user accounts, *almost-theres* and *leads* was comparable between both models. When comparing token usage, their respective PLANNER components used similar amounts of tokens while DEEPSEEK-V3’s EXECUTOR component used roughly half the tokens of GPT-4o. We employed the respective LLM maker’s cloud offerings for hosting the models. Both models generated PTTs that were comparable in size and growth rate (Figure 11). Figure 10(b) shows that DeepSeek’s hosted platform’s response time scales worse compared to OpenAI’s platform. Tool usage was similar for both models. Traces indicate that both models have sufficient penetration-testing background and tool knowledge within their training data set.

5.1.2 Attack Vector Coverage. We used professional penetration-tester to categorize the pursued attack vectors. Covered Attack Vectors converged for each model with similar attack classes covered by both of them (Figure 9). For a qualitative analysis of the covered attack classes see Section 5.1.2.

Both models were able to install missing tools if the used Linux distribution did not include them by default. Both models struggled with successful exploitation: while executed commands were for correct attack vectors and well-executed, the PLANNER was not able to follow up on initial findings. OpenAI’s GPT-4o pursued more attack venues compared to DEEPSEEK-V3. While they did not lead to successful exploitation, GPT-4o’s results included more *almost*